

Problem statement

Adult Salary Prediction

**Solved by:- Iriventi Bharath Vasishta(CSC/20/36),
Mohtassim Rahman(CSC/20/29), Nikhil
Choudhary(CSC/20/42).**

Our task is to analyze the dataset and predict whether the income of an adult will exceed 50k per year or not.

- We have been given 14 features.

Adult dataset link : <http://archive.ics.uci.edu/ml/datasets/Adult>
(<http://archive.ics.uci.edu/ml/datasets/Adult>).

In [110]:

```
import numpy as np # For mathematical calculations
import pandas as pd # For data preprocessing
import matplotlib.pyplot as plt # For data visualization
import seaborn as sns # For data visualization
from sklearn.preprocessing import StandardScaler # For feature scaling
from statsmodels.stats.outliers_influence import variance_inflation_factor # For detect
from sklearn.model_selection import train_test_split # For splitting data into training
from sklearn.metrics import confusion_matrix, classification_report, accuracy_score, pre
from sklearn.tree import DecisionTreeClassifier # For decision tree classification
from sklearn.ensemble import RandomForestClassifier # For random forest classification
from sklearn.svm import SVC # For support vector machine classification
from imblearn.over_sampling import SMOTE # For oversampling the minority class
import warnings
warnings.filterwarnings("ignore") # To ignore warnings for cleaner output
```

In [111]:

```
df = pd.read_csv("adult_data.csv")
df.head()
```

Out[111]:

	age	workclass	education	occupation	race	sex	capital-gain	capital-loss	hours-per-week	native-country
0	39	State-gov	Bachelors	Adm-clerical	White	Male	2174	0	40	United-States
1	50	Self-emp-not-inc	Bachelors	Exec-managerial	White	Male	0	0	13	United-States
2	38	Private	HS-grad	Handlers-cleaners	White	Male	0	0	40	United-States
3	53	Private	11th	Handlers-cleaners	Black	Male	0	0	40	United-States
4	28	Private	Bachelors	Prof-specialty	Black	Female	0	0	40	Cuba

In [112]:

```
df.columns
```

Out[112]:

```
Index(['age', 'workclass', 'education', 'occupation', 'race', 'sex',
      'capital-gain', 'capital-loss', 'hours-per-week', 'native-count
ry',
      'salary'],
      dtype='object')
```

In [113]:

```
len(df.columns)
```

Out[113]:

11

In [114]:

df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                   32561 non-null  int64
1   workclass             32561 non-null  object
2   education             32561 non-null  object
3   occupation            32561 non-null  object
4   race                  32561 non-null  object
5   sex                   32561 non-null  object
6   capital-gain          32561 non-null  int64
7   capital-loss          32561 non-null  int64
8   hours-per-week        32561 non-null  int64
9   native-country        32561 non-null  object
10  salary                32561 non-null  object
dtypes: int64(4), object(7)
memory usage: 2.7+ MB
```

In [115]:

df.shape

Out[115]:

(32561, 11)

In [116]:

df.describe()

Out[116]:


	age	capital-gain	capital-loss	hours-per-week
count	32561.000000	32561.000000	32561.000000	32561.000000
mean	38.581647	1077.648844	87.303830	40.437456
std	13.640433	7385.292085	402.960219	12.347429
min	17.000000	0.000000	0.000000	1.000000
25%	28.000000	0.000000	0.000000	40.000000
50%	37.000000	0.000000	0.000000	40.000000
75%	48.000000	0.000000	0.000000	45.000000
max	90.000000	99999.000000	4356.000000	99.000000

In [117]:

```
sns.countplot(x=' salary',data=df )
```

Out[117]:

<AxesSubplot:xlabel=' salary', ylabel='count'>



In [118]:

```
df[" salary"].value_counts()
```

Out[118]:

```
<=50K    24720  
>50K      7841  
Name: salary, dtype: int64
```

In [119]:

```
print(f"<= 50k : {round(24720 / 32561 * 100 , 2)}%")  
print(f"> 50k : {round(7841 / 32561 * 100 , 2)}%")
```

```
<= 50k : 75.92  
> 50k : 24.08
```

Exploratory data analysis and Feature Engineering

In [120]:

```
## Here we will check the missing values in our dataset
df.isnull().sum()
```

Out[120]:

```
age                0
workclass          0
education          0
occupation        0
race              0
sex               0
capital-gain      0
capital-loss      0
hours-per-week    0
native-country    0
salary            0
dtype: int64
```

In [121]:

```
numerical = [i for i in df.columns if df[i].dtypes != "O"]
```

In [122]:

```
numerical
```

Out[122]:

```
['age', ' capital-gain', ' capital-loss', ' hours-per-week']
```


In [123]:

```
for i in numerical:
    print(f" {i} : {len(df[i].unique())}")
```

```
age : 73
capital-gain : 119
capital-loss : 92
hours-per-week : 94
```

In [124]:

```
for feature in numerical:
    bar = sns.distplot(df[feature] , kde_kws = {'bw' : 1})
    bar.legend(["Skewness: {:.2f}".format(df[feature].skew())])
    plt.xlabel(feature)
    plt.ylabel("Probability density")
    plt.title(feature)
    plt.show()
```



In [125]:

```
df["age"] = np.log(df["age"])
```

In [126]:

```
df.columns
```

Out[126]:

```
Index(['age', 'workclass', 'education', 'occupation', 'race', 'sex',
      'capital-gain', 'capital-loss', 'hours-per-week', 'native-count
ry',
      'salary'],
      dtype='object')
```

Handling categorical data

In [127]:

```
categorical = [i for i in df.columns if df[i].dtypes == "O"]
```

In [128]:

```
for feature in categorical:
    print(f" {feature} : {len(df[feature].unique())}")

workclass : 9
education : 16
occupation : 15
race : 5
sex : 2
native-country : 42
salary : 2
```

In [129]:

```
for feature in categorical:
    print(df[feature].value_counts())
    print("\n \n ")
```

```
Cambodia 19
Trinidad&Tobago 19
Laos 18
Thailand 18
Yugoslavia 16
Outlying-US(Guam-USVI-etc) 14
Honduras 13
Hungary 13
Scotland 12
Holand-Netherlands 1
```

Name: native-country, dtype: int64

```
<=50K 24720
>50K 7841
```

Name: salary, dtype: int64

Handling missing values in categorical features

In [130]:

```
print(f"workclass : {round(2093 / 32561 , 4) *100}%")
print(f"occupation : {round(1843 / 32561 , 4) *100}%")
print(f"native-country : {round(583 / 32561 , 4) *100}%")
```

```
workclass : 6.43%
occupation : 5.66%
native-country : 1.79%
```

In [131]:

```
df[" occupation"].mode()[0]
```

Out[131]:

```
' Prof-specialty'
```

here we have less than 6 percent missing values so we can fill it with mode value

In [132]:

```
df[" workclass"] = df[' workclass'].str.replace('?', 'Private' )
df[' occupation'] = df[' occupation'].str.replace('?', 'Prof-specialty' )
df[' native-country'] = df[' native-country'].str.replace('?', 'United-States' )
```

In [133]:

```
# education Category
df[" education"].replace(['Preschool', '1st-4th', '5th-6th', '7th-8th', '9th','10th', '11th', '12th'],
                        inplace = True , regex = True)
df[" education"].replace(['Assoc-voc', 'Assoc-acdm', 'Prof-school', 'Some-college'], 'high-school
```

In [134]:

```
# income
df[" salary"] = df[" salary"].replace({'<=50K' : 0 , ">50K" : 1 } , regex = True)
```

In [135]:

```
df.head()
```

Out[135]:

	age	workclass	education	occupation	race	sex	capital-gain	capital-loss	hours-per-week	native-country
0	3.663562	State-gov	Bachelors	Adm-clerical	White	Male	2174	0	40	United-States
1	3.912023	Self-emp-not-inc	Bachelors	Exec-managerial	White	Male	0	0	13	United-States
2	3.637586	Private	HS-grad	Handlers-cleaners	White	Male	0	0	40	United-States
3	3.970292	Private	school	Handlers-cleaners	Black	Male	0	0	40	United-States
4	3.332205	Private	Bachelors	Prof-specialty	Black	Female	0	0	40	Canada

In [136]:

```
for feature in categorical:
    print(f" {feature} : {len(df[feature].unique())}")
```

```
workclass : 8
education : 6
occupation : 14
race : 5
sex : 2
native-country : 41
salary : 2
```


In [137]:


```
df[" education"].unique()
```

Out[137]:

```
array([' Bachelors', ' HS-grad', ' school', ' Masters', ' higher',  
      ' Doctorate'], dtype=object)
```

In [138]:

```
plt.figure(figsize = (16 , 7))  
sns.heatmap(df.corr(), annot=True);
```



In [139]:

```
from sklearn.preprocessing import LabelEncoder
```

In [140]:

```
df = df.apply(LabelEncoder().fit_transform)  
df .head()
```

Out[140]:

	age	workclass	education	occupation	race	sex	capital-gain	capital-loss	hours-per-week	native-country	sala
0	22	6	0	0	4	1	25	0	39	38	
1	33	5	0	3	4	1	0	0	12	38	
2	21	3	2	5	4	1	0	0	39	38	
3	36	3	5	5	2	1	0	0	39	38	
4	11	3	0	9	2	0	0	0	39	4	



Pre Processing

In [141]:

```
# Preprocess the data for machine Learning
target_name = ' salary'
y = df[target_name]
x_train = df.drop(target_name, axis=1)
```

In [142]:

```
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
x_res = sc.fit_transform(x_train)
```

In [143]:

```
from sklearn.metrics import confusion_matrix, classification_report, accuracy_score
```

In [144]:

```
# Calculate variance inflation factor (VIF) for each feature
vif = pd.DataFrame()
vif["VIF Factor"] = [variance_inflation_factor(x_res, i) for i in range(x_res.shape[1])]
vif["features"] = x_train.columns
```

In [145]:

```
# Display VIF for each feature
print("Variance Inflation Factor (VIF):")
pd.DataFrame(vif)
```

Variance Inflation Factor (VIF):

Out[145]:

	VIF Factor	features
0	1.031483	age
1	1.008832	workclass
2	1.024136	education
3	1.003528	occupation
4	1.024756	race
5	1.079126	sex
6	1.036486	capital-gain
7	1.015167	capital-loss
8	1.084125	hours-per-week
9	1.016311	native-country

In [146]:

```
x_res.shape
```

Out[146]:

```
(32561, 10)
```

MODEL EVALUATION

1) Decision Tree Classifier

In [147]:

```
X_train, X_test, y_train, y_test = train_test_split(x_res, y, test_size=0.2, random_stat

# Print the shapes of training and testing sets
print("X_train.shape:", X_train.shape)
print("y_train.shape:", y_train.shape)
print("X_test.shape:", X_test.shape)
print("y_test.shape:", y_test.shape)
```

```
X_train.shape: (26048, 10)
y_train.shape: (26048,)
X_test.shape: (6513, 10)
y_test.shape: (6513,)
```

In [148]:

```
dt = DecisionTreeClassifier()
dt.fit(X_train, y_train)

# Predict on training and testing sets using Decision Tree Classifier
dt_train_pred = dt.predict(X_train)
dt_test_pred = dt.predict(X_test)
```

In [149]:

```
# Print accuracy scores and confusion matrix for Decision Tree Classifier
print('Train Accuracy:', accuracy_score(y_train, dt_train_pred) * 100)
print('Accuracy Score:', accuracy_score(y_test, dt_test_pred) * 100)
```

```
Train Accuracy: 95.98817567567568
Accuracy Score: 79.57930293259633
```

Confusion Matrix

In [150]:

```
print(confusion_matrix(y_test, dt_test_pred))
```

```
[[4330  621]
 [ 709  853]]
```

In [151]:

```
#Classification_report
```

In [152]:

```
print(classification_report(y_test, dt_test_pred, digits=4))
```

	precision	recall	f1-score	support
0	0.8593	0.8746	0.8669	4951
1	0.5787	0.5461	0.5619	1562
accuracy			0.7958	6513
macro avg	0.7190	0.7103	0.7144	6513
weighted avg	0.7920	0.7958	0.7937	6513

Precision Score

In [153]:

```
print('Precision score of macro is: ', round(precision_score(y_test, dt_test_pred, av
print('Precision score of micro is: ', round(precision_score(y_test, dt_test_pred, av
print('Precision score of weighted is: ', round(precision_score(y_test, dt_test_pred, av
```

```
Precision score of macro is: 71.9
Precision score of micro is: 79.58
Precision score of weighted is: 79.2
```

Recall Score

In [154]:

```
print('recall_score score of macro is: ', round(recall_score(y_test, dt_test_pred, av
print('recall_score score of micro is: ', round(recall_score(y_test, dt_test_pred, av
print('recall_score score of weighted is: ', round(recall_score(y_test, dt_test_pred, av
```

```
recall_score score of macro is: 71.03
recall_score score of micro is: 79.58
recall_score score of weighted is: 79.58
```

F1 Score

In [155]:

```
print('f1_score score of macro is: ', round(f1_score(y_test, dt_test_pred, average =
print('f1_score score of micro is: ', round(f1_score(y_test, dt_test_pred, average =
print('f1_score score of weighted is: ', round(f1_score(y_test, dt_test_pred, average
```

```
f1_score score of macro is: 71.44
f1_score score of micro is: 79.58
f1_score score of weighted is: 79.37
```

2)Random Forest

In [156]:

```
from sklearn.ensemble import RandomForestClassifier
rf_classifier = RandomForestClassifier(n_estimators = 20, criterion = 'entropy', random_
rf_classifier.fit(X_train, y_train)
y_pred = rf_classifier.predict(X_test)
accuracy_score(y_test, y_pred)*100
```

Out[156]:

82.84968524489483

In [157]:


```
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.87	0.92	0.89	4951
1	0.68	0.55	0.60	1562
accuracy			0.83	6513
macro avg	0.77	0.73	0.75	6513
weighted avg	0.82	0.83	0.82	6513

Confusion Matrix

In [158]:

```
cm = confusion_matrix(y_test, y_pred )
plt.title('Heatmap of Confusion Matrix', fontsize = 12)
sns.heatmap(cm, annot = True , fmt = "d")
plt.show()
```



3)Support Vector Machine(SVM)

In [159]:

```
# Train Support Vector Classifier (SVC)
svc_model = SVC()
svc_model.fit(X_train, y_train)

# Predict on training and testing sets using SVC
svc_train_pred = svc_model.predict(X_train)
svc_test_pred = svc_model.predict(X_test)
```

In [160]:

```
# Print accuracy scores and confusion matrix for SVC
print('Train Accuracy:', accuracy_score(y_train, svc_train_pred) * 100)
print('Accuracy Score:', accuracy_score(y_test, svc_test_pred) * 100)
```

Train Accuracy: 82.8585687960688
 Accuracy Score: 82.081989866421

Confusion Matrix

In [161]:

```
print(confusion_matrix(y_test, svc_test_pred))
```

```
[[4777  174]
 [ 993  569]]
```

Classification Report

In [66]:

```
print(classification_report(y_test, svc_test_pred, digits=4))
```

	precision	recall	f1-score	support
0	0.8279	0.9649	0.8911	4951
1	0.7658	0.3643	0.4937	1562
accuracy			0.8208	6513
macro avg	0.7969	0.6646	0.6924	6513
weighted avg	0.8130	0.8208	0.7958	6513

Handling Imbalanced data

In [162]:

```
# Check value counts of y
print(y.value_counts())
```

```
0    24720
1     7841
Name: salary, dtype: int64
```

In [163]:

```
# Apply SMOTE for handling imbalanced dataset
oversample = SMOTE(k_neighbors=4)
X_res1, y_res1 = oversample.fit_resample(x_res, y)
```

In [164]:

```
# Print shapes of resampled dataset
print("X_res1.shape:", X_res1.shape)
print("y_res1.shape:", y_res1.shape)
```

```
X_res1.shape: (49440, 10)
y_res1.shape: (49440,)
```

In [165]:

```
# Check value counts of original and resampled y
print("Original y value counts:")
print(y.value_counts())
print("Resampled y value counts:")
print(pd.Series(y_res1).value_counts())
```

Original y value counts:

0 24720

1 7841

Name: salary, dtype: int64

Resampled y value counts:


0 24720

1 24720

Name: salary, dtype: int64

In [166]:

```
# Plot bar chart of resampled y value counts
pd.Series(y_res1).value_counts().plot.bar()
plt.grid()
```



Retraining and Re-evaluation

1) Decision Tree Classifier

In [167]:

```
X_train, X_test, y_train, y_test = train_test_split(X_res1, y_res1, test_size=0.2, random_state=42)

print("X_train.shape:", X_train.shape)
print("y_train.shape:", y_train.shape)
print("X_test.shape:", X_test.shape)
print("y_test.shape:", y_test.shape)
```

```
X_train.shape: (39552, 10)
y_train.shape: (39552,)
X_test.shape: (9888, 10)
y_test.shape: (9888,)
```

In [168]:

```
# Decision Tree Classifier
dt = DecisionTreeClassifier()
dt.fit(X_train, y_train)

dt_train_pred = dt.predict(X_train)
dt_test_pred = dt.predict(X_test)
```

In [169]:

```
print('Decision Tree Classifier (after SMOTE) Train Accuracy: ', accuracy_score(y_train, dt_train_pred))
print('Decision Tree Classifier (after SMOTE) Accuracy Score: ', accuracy_score(y_test, dt_test_pred))
```

```
Decision Tree Classifier (after SMOTE) Train Accuracy: 96.93315129449837
Decision Tree Classifier (after SMOTE) Accuracy Score: 83.59627831715211
```

Confusion Matrix

In [170]:

```
print(confusion_matrix(y_test, dt_test_pred))
```

```
[[4224  723]
 [ 899 4042]]
```

Classification Report

In [171]:

```
print(classification_report(y_test, dt_test_pred, digits=4))
```

	precision	recall	f1-score	support
0	0.8245	0.8539	0.8389	4947
1	0.8483	0.8181	0.8329	4941
accuracy			0.8360	9888
macro avg	0.8364	0.8360	0.8359	9888
weighted avg	0.8364	0.8360	0.8359	9888

2) Random Forest

In [172]:

```
# Random Forest Classifier
from sklearn.ensemble import RandomForestClassifier
rf_classifier = RandomForestClassifier(n_estimators = 20, criterion = 'entropy', random_
rf_classifier.fit(X_train, y_train)
y_pred = rf_classifier.predict(X_test)
accuracy_score(y_test, y_pred)*100
```

Out[172]:

86.5796925566343

Classification Report

In [173]:


```
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.86	0.87	0.87	4947
1	0.87	0.86	0.86	4941
accuracy			0.87	9888
macro avg	0.87	0.87	0.87	9888
weighted avg	0.87	0.87	0.87	9888

Confusion Report

In [174]:

```
cm = confusion_matrix(y_test, y_pred )
plt.title('Heatmap of Confusion Matrix', fontsize = 12)
sns.heatmap(cm, annot = True , fmt = "d")
plt.show()
```



3)SVM

In [175]:

```
# Train Support Vector Classifier (SVC)
svc_model = SVC()
svc_model.fit(X_train, y_train)

# Predict on training and testing sets using SVC
svc_train_pred = svc_model.predict(X_train)
svc_test_pred = svc_model.predict(X_test)
```

In [176]:

```
# Print accuracy scores and confusion matrix for SVC
print('Train Accuracy:', accuracy_score(y_train, svc_train_pred) * 100)
print('Accuracy Score:', accuracy_score(y_test, svc_test_pred) * 100)
```

Train Accuracy: 78.85821197411003
 Accuracy Score: 78.47896440129449

Confusion Report

In [177]:

```
print(confusion_matrix(y_test, svc_test_pred))
```

```
[[3752 1195]
 [ 933 4008]]
```

Classification Report

In [178]:

```
print(classification_report(y_test, svc_test_pred, digits=4))
```

	precision	recall	f1-score	support
0	0.8009	0.7584	0.7791	4947
1	0.7703	0.8112	0.7902	4941
accuracy			0.7848	9888
macro avg	0.7856	0.7848	0.7846	9888
weighted avg	0.7856	0.7848	0.7846	9888

In []: